

# Week 1: Introduction and overview

Text Analytics and Natural Language Processing  
Instructor: Benjamin Batorsky

# Explosion of data...unstructured data, that is

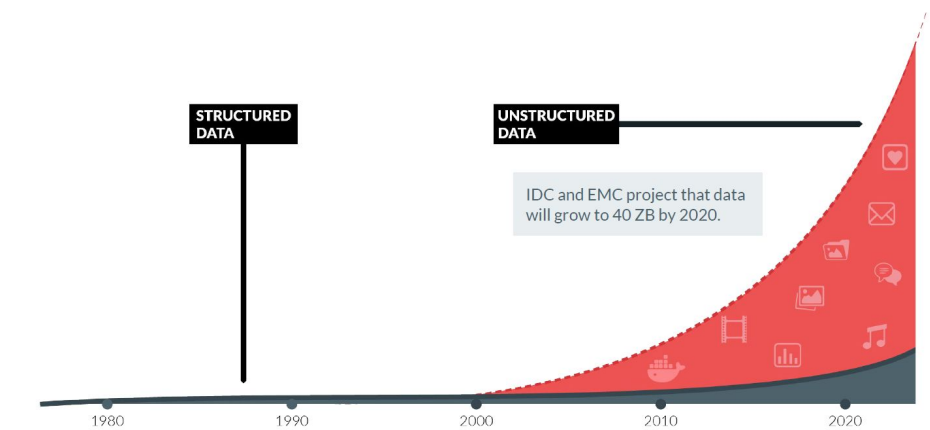

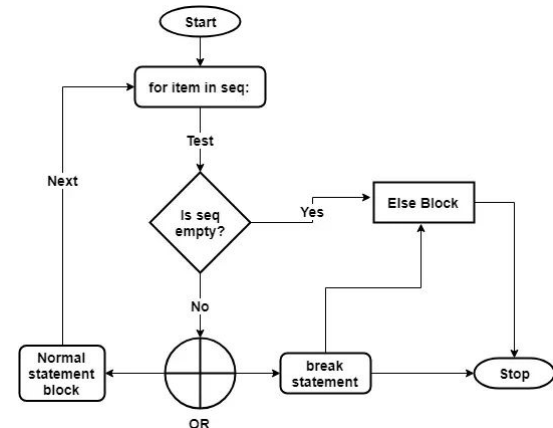



Chart: <https://www.datanami.com/2017/02/01/solving-storage-just-beginning-minio-ceo-periasamy/>  
Data: IDC Structured Versus Unstructured Data: The Balance of Power Continues to Shift, March 2014

# What is Natural Language?

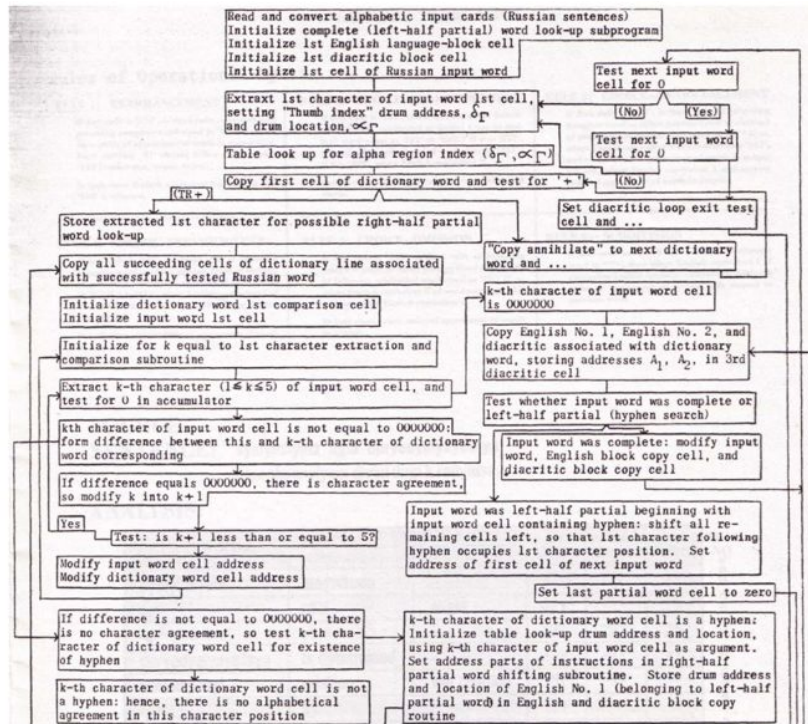
*"A language that has developed naturally in use (as contrasted with an artificial language or computer code)."*  
(Oxford Dictionary definition)



<https://www.techbeamers.com/python-for-loop/>

# How do we process language into text?

1950



2013



<https://jalamar.github.io/visualizing-neural-machine-translation-mechanics-of-seq2seq-models-with-attention/>

# Now we can do things like this

Write with Transformer from HuggingFace:

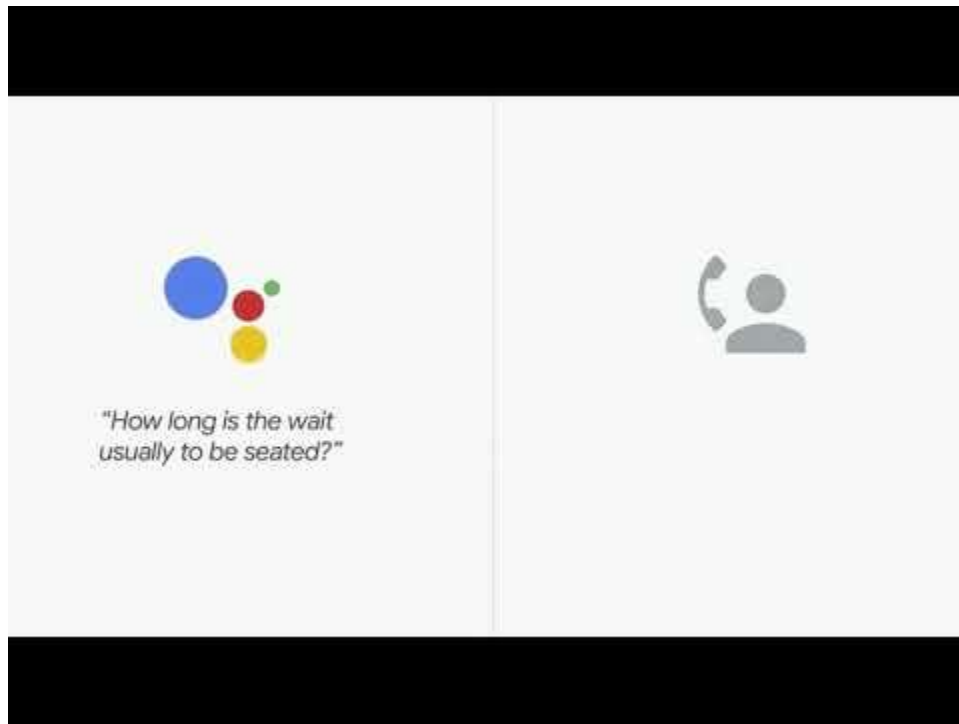
<https://transformer.huggingface.co/doc/distil-gpt2>

**I am teaching a course at Harvard on how to develop and learn about language skills,  
such as German and French.**

Written by Transformer · [transformer.huggingface.co](https://transformer.huggingface.co) 🦄


# And this:

Google Assistant making a reservation



# Though also, this:

Word vector similarity between gender words and occupations



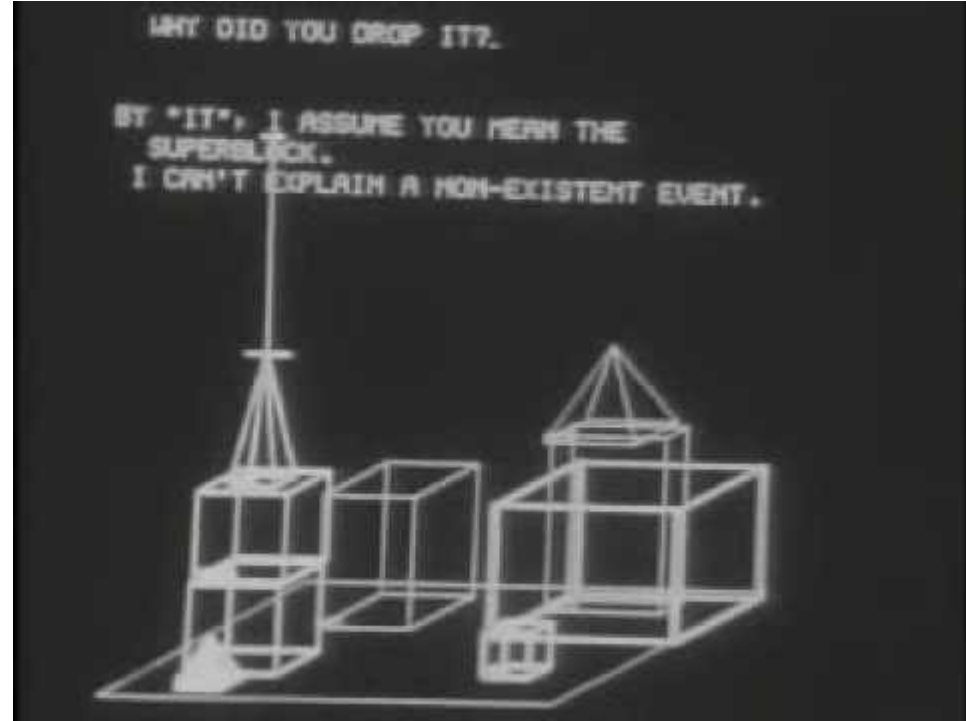
[Exploring and Mitigating Gender Bias in GloVe Word Embeddings](#)





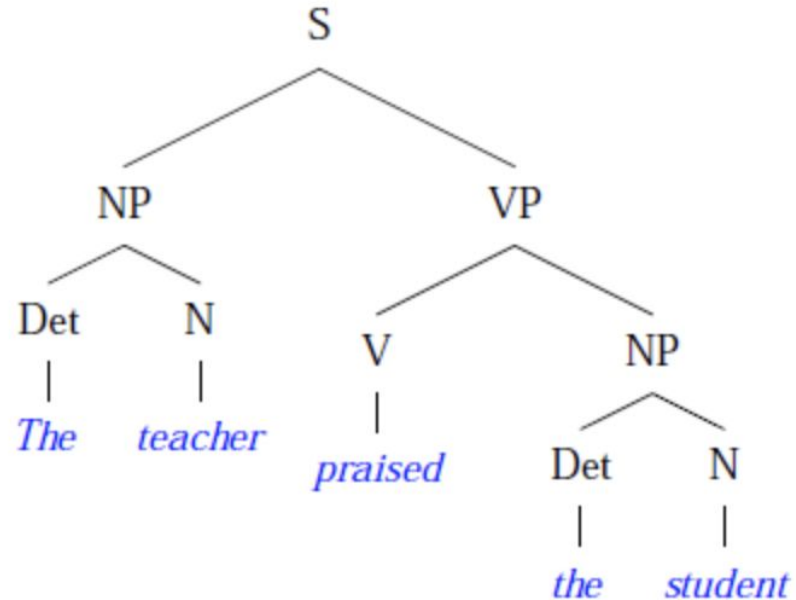
# 1960s-1970s: Shift to semantic-driven processing

- Semantic-driven processing
  - Focus on meaning/relationships
  - Incorporate information outside of the language itself
- SHRLDU
  - Map query to stored relationship map of the machine state
- Focus on pulling in different data sources
  - Lots of focus on tools for data integration
- Still not enough
  - Uninteresting results on actual dealing with natural language
  - User-machine exchanges difficult to handle “naturally”




# 1970s-1980s: Joining of linguistics and AI

- Linguistics (Study of language)
  - Noam Chomsky's ideas of "generative grammar"
  - Focus on categorization/computability
- AI
  - Increasing focus on representations of meanings/intention
- Joining together, building of community
  - New conferences
  - New standardized tasks to measure performance



# 1990s-2000s: Shift towards statistical NLP

- Increasing attention on probabilistic models of language
  - Visual of parsing trees with probabilities, Markov Models
  - Probabilities of certain words appearing, the precursor to modern language models
- Also newly available data sources for developing/testing
  - WordNet (<https://wordnet.princeton.edu/>)
    - Inventory of English words and their relationships
  - Penn TreeBank (<https://catalog.ldc.upenn.edu/LDC99T42>)
    - Text parsed for part of speech, dependencies, etc
  - Ontonotes (<https://catalog.ldc.upenn.edu/LDC2013T19>)
    - Text labelled with entities (e.g. places and persons)
  - The internet!



# To the notebooks!

[https://github.com/bpben/nlp\\_lessons/blob/master/notebooks\\_instructor/week\\_1\\_intro.ipynb](https://github.com/bpben/nlp_lessons/blob/master/notebooks_instructor/week_1_intro.ipynb)

# NLP libraries

Comparison of Python NLP libraries Pros and Cons		
	☀ PROS	☹ CONS
	<ul style="list-style-type: none"> <li>+ The most well-known and full NLP library</li> <li>+ Many third-party extensions</li> <li>+ Plenty of approaches to each NLP task</li> <li>+ Fast sentence tokenization</li> <li>+ Supports the largest number of languages compared to other libraries</li> </ul>	<ul style="list-style-type: none"> <li>- Complicated to learn and use</li> <li>- Quite slow</li> <li>- In sentence tokenization, NLTK only splits text by sentences, without analyzing the semantic structure</li> <li>- Processes strings which is not very typical for object-oriented language Python</li> <li>- Doesn't provide neural network models</li> <li>- No integrated word vectors</li> </ul>
	<ul style="list-style-type: none"> <li>+ The fastest NLP framework</li> <li>+ Easy to learn and use because it has one single highly optimized tool for each task</li> <li>+ Processes objects; more object-oriented, comparing to other libs</li> <li>+ Uses neural networks for training some models</li> <li>+ Provides built-in word vectors</li> <li>+ Active support and development</li> </ul>	<ul style="list-style-type: none"> <li>- Lacks flexibility, comparing to NLTK</li> <li>- Sentence tokenization is slower than in NLTK</li> <li>- Doesn't support many languages. There are models only for 7 languages and "multi-language" models</li> </ul>
	<ul style="list-style-type: none"> <li>+ Has functions which help to use the bag-of-words method of creating features for the text classification problems</li> <li>+ Provides a wide variety of algorithms to build machine learning models</li> <li>+ Has good documentation and intuitive classes' methods</li> </ul>	<ul style="list-style-type: none"> <li>- For more sophisticated preprocessing things (for example, pos-tagging), you should use some other NLP library and only after it you can use models from scikit-learn</li> <li>- Doesn't use neural networks for text preprocessing</li> </ul>
	<ul style="list-style-type: none"> <li>+ Works with large datasets and processes data streams</li> <li>+ Provides tf-idf vectorization, word2vec, document2vec, latent semantic analysis, latent Dirichlet allocation</li> <li>+ Supports deep learning</li> </ul>	<ul style="list-style-type: none"> <li>- Designed primarily for unsupervised text modeling</li> <li>- Doesn't have enough tools to provide full NLP pipeline, so should be used with some other library (Spacy or NLTK)</li> </ul>
	<ul style="list-style-type: none"> <li>+ Allows part-of-speech tagging, n-gram search, sentiment analysis, WordNet, vector space model, clustering and SVM</li> <li>+ There are web crawler, DOM parser, some APIs (like Twitter, Facebook etc.)</li> </ul>	<ul style="list-style-type: none"> <li>- Is a web miner; can be not enough optimized for some specific NLP tasks</li> </ul>
	<ul style="list-style-type: none"> <li>+ Supports a large number of languages (16-196 languages for different tasks)</li> </ul>	<ul style="list-style-type: none"> <li>- Not as popular as, for example, NLTK or Spacy; can be slow issues solutions or weak community support</li> </ul>

[Comparison of Top 6 Python NLP Libraries - ActiveWizards — your AI partner](#)

# Note on Stanza (formerly StanfordNLP)

- <https://github.com/stanfordnlp/stanza/>
- Similar “pipeline” structure to spaCy
  - `stanza.download('en')`
  - `nlp = stanza.Pipeline('en')`
- Large amount of languages models (POS/NER)
  - [https://stanfordnlp.github.io/stanza/available\\_models.html](https://stanfordnlp.github.io/stanza/available_models.html)
- Spacy-Stanza
  - <https://spacy.io/universe/project/spacy-stanza>
  - Most of the same attributes available
- If looking for a language that's not available consider Stanza or <http://ufal.mff.cuni.cz/udpipe>

# Tokenization

- Definition: “A useful semantic unit”
- What is a “useful semantic unit”?

“I am learning Natural Language Processing (NLP)”

What are the useful units of meaning here?

I, am, learning (split on whitespace)

Natural Language Processing (entity)

NLP (entity as acronym)

(NLP) (special treatment of text in parentheses?)

# N-grams

- N-gram: Continuous sequence of N tokens
  - Tokens != words (see previous slides)
- Important considerations
  - “not”, “good” vs “not good”
  - Named-entities (e.g. Cities, people)
  - Exploring co-occurrence

$\text{count}(\text{natural language}) / \text{count}(\text{language})$

$\text{count}(\text{english language}) / \text{count}(\text{language})$

“I am learning Natural Language Processing  
(NLP)”

<split on whitespace>

Unigrams

I, am, learning, Natural, Language, Processing,  
(NLP)

Bigrams

I am, am learning, learning Natural...

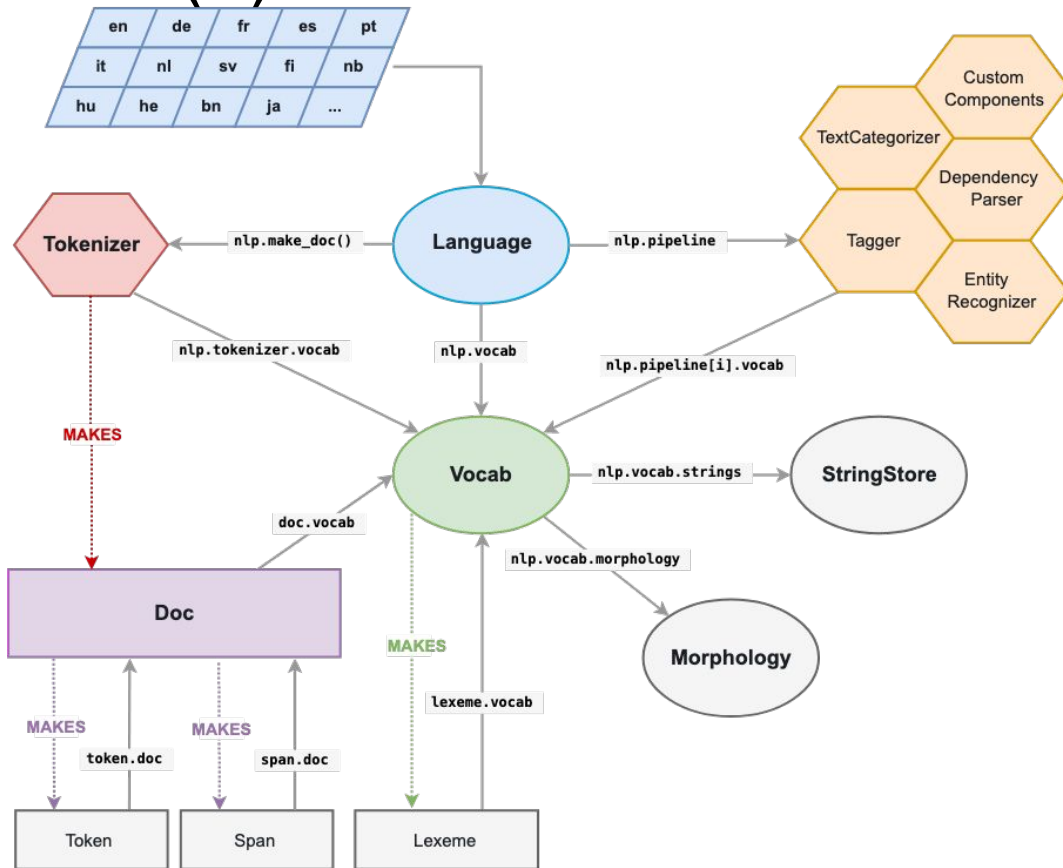
7-grams

I am learning Natural Language Processing  
(NLP)




# SpaCy's Language models (1)

- Language model:  
Tokenization/processing pipeline
- Base: Tokenizer + Vocab
  - Tokenizer
    - Processing of text
  - Vocab
    - Contains all “lexemes”
- Document -> Span -> Token
  - .text = raw text
  - Document/Span
  - Token
    - .shape\_
      - We = Xx
      - we = xx
    - .is\_alpha (boolean)
    - .is\_stop (boolean)



# SpaCy's Language models (2)

- Can load “trained” language models
- Additional pipeline components
  - Entity recognition
  - Part of speech tagging
  - Dependency parsing
  - Categorization
  - Anything else (custom attributes)
- Products of these
  - Document/Span
    - .ents (named entities)
    - .sents
    - .noun\_chunks
  - Token
    - .dep\_ (e.g. nsubj)
    - .lemma\_ (dictionary form)
    - .pos\_ (part of speech)



# What is the point of NLP?

What is the main challenge of working with text data?

How is it different from structured data (e.g. stock ticker, height/weight)?

What are we looking for when we're trying to find these “useful semantic units”?

**Main goal: Create an informative representation of text**

# Stemming vs Lemmatization

Goal: Reduce related words to a common “base form”

## Stemming

Set of heuristics to transform suffixes. Most of the word stays intact.

“He does natural language processing”



“He **doe** natural language **process**”

Implemented in: NLTK, PyStemmer

## Lemmatization

Transforms word to their “lemma”, or dictionary form. This may change the word entirely.

“He is doing natural language processing”



“He **be** **do** natural language **process**”

Implemented in: NLTK, spaCy

# Stop words

- Extremely common words that have little information
  - “the”, “at”, “from”, “to”, “in”
- No single list of these
  - Consequences to choosing overly strict/broad

**In June 2020, I took a course  
at Harvard Extension School  
in Cambridge.**

<Remove stop words (SpaCy)>

**June 2020 took course  
Harvard Extension School  
Cambridge**

# Issues with stop words

Scikit-learn

Example: “hasn’t” is removed using “hasnt”

[https://github.com/scikit-learn/scikit-learn/blob/fd237278e895b42abe8d8d09105cbb82dc2cbba7/sklearn/feature\\_extraction/\\_stop\\_words.py](https://github.com/scikit-learn/scikit-learn/blob/fd237278e895b42abe8d8d09105cbb82dc2cbba7/sklearn/feature_extraction/_stop_words.py)

SpaCy

Example: “hasn’t” is removed using “has” and contraction “n’t”

[https://github.com/explosion/spaCy/blob/master/spacy/lang/en/stop\\_words.py](https://github.com/explosion/spaCy/blob/master/spacy/lang/en/stop_words.py)




Figure 1: Family tree of popular stop word lists.

<https://www.aclweb.org/anthology/W18-2502.pdf>

# Named-Entities

- Named-entity: A real-world named object (e.g. person, place, organization)
  - New York City is different than just an assembly of three words “new”, “york” and “city”
- To identify these
  - Dictionaries
  - Pattern-matching
  - Models

In **June 2020 [DATE]**, I took a course at **Harvard Extension School [ORG]** in **Cambridge [LOC]**.

# Exercise: Create a tokenization pipeline

[https://github.com/bpben/nlp\\_lessons/blob/master/notebooks\\_instructor/week\\_1\\_intro.ipynb](https://github.com/bpben/nlp_lessons/blob/master/notebooks_instructor/week_1_intro.ipynb)



# Log-likelihood ratio with word counts

- Enables comparison of word-level counts between two documents/corpora
- Test of hypothesis: Frequency of word is equivalent between analysis and reference
  - Expected count = number of words in document  $i$  \* the pooled frequency across both documents
- Provides a test of significance for differences in frequency
- Accessible via simple word counts!

	Analysis Text	Reference Text	Total
Count of word form	a	b	a+b
Count of other word forms	c-a	d-b	c+d-a-b
Total	c	d	c+d

$$E1 = c * (a+b) / (c+d)$$


$$E2 = d * (a+b) / (c+d)$$

$$G^2 = 2 * ((a * \ln(a/E1)) + (b * \ln(b/E2)))$$

<https://wordhoard.northwestern.edu/userman/analysis-comparewords.html>


# 2000s: Advent of Neural Models for NLP

- Focused on language modelling task
  - Given context words, predict next word
  - “Do you want to go out for <mask>”
  - “I ate so much I am so <mask>”
  - “I’m so tired, I think I’ll take a <mask>”
- Apply Neural Net architectures to language modelling task
  - Neural Net: Model that connects inputs to outputs through sets of computations
- Bengio et. al. (2001) [A Neural Probabilistic Language Model](#)
  - Context words fed into a matrix that “represents” the information
  - These representations then fed into a computation layer
  - Output: Prediction of the target word
    - “Full”: 70% likely, “tired” 20% likely, etc




# 2008: Multi-task learning

- Multi-task learning: One model, multiple related tasks
  - Trains representations (e.g. lookup tables) jointly so that it performs well on both
- Collobert and Weston (2008)
  - Word lookup table trained jointly from two different tasks
  - Basically the precursor to the idea of word embeddings




# 2013: Word embeddings


- Mikolov paper
  - Efficient computation of embedding technique used in the 2000s
  - Word2Vec implementation
- Huge amount of adaptation/modifications
  - GloVe: Co-occurrence based
  - Sub-word/character-level
  - Use of different corpora for training
  - Out-of-vocabulary handling
  - Training implementation in open-source libraries (e.g. gensim)
- Major revolution: It's fairly easy to make more informed representations of words
- Caveat: Widespread use raises issues of bias
  - Extensive research finds training on most large datasets introduces gender/racial bias
  - Not an issue with the method particularly, more with irresponsible use



Male-Female



Verb tense



Country-Capital

[\[1301.3781\] Efficient Estimation of Word Representations in Vector Space](#)